Yunfeng Zhang

Machine Learning Engineer / AI Researcher

Chappaqua, NY 10514
zywind@gmail.com
linkedin.com/in/zywind/

SUMMARY

- AI/HCI researcher dedicated to developing fair and interpretable AI systems.
- Proficient in machine learning operations (MLOps) and infrastructure development.
- Published over 50 papers with more than 2000 citations. Filed 13 patent applications.

EXPERIENCE

Twitter Inc., NYC - *Senior Machine Learning Engineer*

June 2021 - PRESENT

- Architected and developed Fairness Evaluator, Twitter's first AI fairness evaluation system based on Tensorflow and Apache Beam. Fairness Evaluator has been used throughout the company to evaluate and improve recommendation, ranking, and content moderation models and is a cornerstone of responsible machine learning at Twitter.
- Researched and developed new AI fairness and performance metrics, including the metrics published at ACM FAccT 2022 (citation below).

IBM, T. J. Watson Research Center - *Research Staff Member*

June 2015 - May 2021

- Developed IBM's **AI Fairness 360** and **AI Explainability 360** toolkits, one of the first open source ethical AI toolkits in the industry. They were the foundation of the IBM OpenScale ML cloud service. Received **IBM outstanding research accomplishment award**.
- Researched and developed various prototype systems on active learning, drift detection, and chatbot authoring to inform designs of production systems such as Watson OpenScale, IBM AutoAI, and IBM Watson Assistant. Received **IBM outstanding accomplishment award**.
- Designed and implemented AI-driven multimodal interaction techniques for smart meeting rooms by incorporating gesture, speech, and face recognition techniques.

IBM, T. J. Watson Research Center - *Research Intern*

May 2014 - September 2014

- Researched methods to remediate cognitive biases in AI-assisted human decision making.

Palo Alto Research Center - *Research Intern*

May 2013 - December 2013

- Developed computational models to simulate and predict how humans detect changes in stochastic environments.

EDUCATION

University of Oregon, Eugene - *Ph.D. in Computer and Information Science,* June 2015

University of Oregon, Eugene - *M.S. in Computer and Information Science,* June 2013

Beijing Normal University, Beijing - *B.S. in Computer and Information Science,* June 2007

PUBLICATIONS and PATENTS

Published over 50 papers with more than 2000 citations, 3 patents, and 9 pending patent applications. See Google scholar page for more information . Selected publications:

- De-biasing "bias" measurement
- AI Fairness 360: An extensible toolkit for detecting, understanding, and mitigating unwanted algorithmic bias
- One explanation does not fit all: A toolkit and taxonomy of AI explainability techniques

AWARDS

- IBM Outstanding Accomplishment Award for Research Advancements to Conversational Technology, 2020.
- IBM Outstanding Research Accomplishment Award for Trustworthy AI, 2019.
- Annual Conference of the Cognitive Science Society, Computational Modeling Award for Applied Cognition, 2014.
- ACM CHI Conference on Human Factors in Computing Systems, Best Paper award, 2014.
- First place, Green Driver Programming Contest, 2011.
- First place, Fifth Annual UO Eugene Luks Programming Contest, 2011.
- ACM CHI Conference on Human Factors in Computing Systems, Honorable Mention award, 2010.
- International Conference on Cognitive Modeling, Siegel-Wolf Award for Best Applied Paper, 2010.

TECHNICAL SKILLS

- Proficient in Python, Java, and R. Familiar with C++, Scala, and Julia.
- Proficient in Tensorflow, Keras, PyTorch, scikit-learn, pandas, numpy.
- Proficient in Apache Beam, TFX, Kubeflow, Apache Spark, GCP, BigQuery, Dataflow.